

Summer 2020

Unix/Linux for Informatic Analysis

https://github.com/sabrsyed/InformaticsTools_2020

Week 1


- Unix Primer for Biologists: Chapters U1 – U16
 - Learn how to use UNIX/Linux
- Logging in to the Cluster
 - Learn to navigate the Cluster
- Powerpoint Presentation: learn the technology behind genome sequencing, what does ChIP-Seq data look like
- Pipeline for ChIP alignment

Week 2

- Unix Primer for Biologists: Chapters U17 – U34
 - Learn how to use UNIX/Linux
- Filezilla
 - Uploading files to the Cluster
- Powerpoint Presentation: RNA-Sequencing, what does RNA-Seq data look like
- Pipeline for RNA-Seq alignment

Optional Exercise:

https://github.com/sabrsyed/InformaticsTools_2020/blob/master/01_Unix_QuickReview_ProblemSet.md



Summer 2020

Unix/Linux for Informatic Analysis

https://github.com/sabrsyed/InformaticsTools_2020

Week 1

- Unix Primer for Biologists: Chapters U1 – U16
 - Learn how to use UNIX/Linux
- Logging in to the Cluster
 - Learn to navigate the Cluster
- Powerpoint Presentation: learn the technology behind genome sequencing, what does ChIP-Seq data look like
- Pipeline for ChIP alignment

Week 2

- Unix Primer for Biologists: Chapters U17 – U34
 - Learn how to use UNIX/Linux
- Filezilla
 - Uploading files to the Cluster
- Powerpoint Presentation: RNA-Sequencing, what does RNA-Seq data look like
- Pipeline for RNA-Seq alignment

Optional Exercise:

https://github.com/sabrsyed/InformaticsTools_2020/blob/master/01_Unix_QuickReview_ProblemSet.md

Unix Primer for Biologists, Chapters U1-U16

Summer 2020

Unix/Linux for Informatic Analysis

https://github.com/sabrsyed/InformaticsTools_2020

Week 1

- Unix Primer for Biologists: Chapters U1 – U16
 - Learn how to use UNIX/Linux
- Logging in to the Cluster
 - Learn to navigate the Cluster
- Powerpoint Presentation: learn the technology behind genome sequencing, what does ChIP-Seq data look like
- Pipeline for ChIP alignment

Week 2

- Unix Primer for Biologists: Chapters U17 – U34
 - Learn how to use UNIX/Linux
- Filezilla
 - Uploading files to the Cluster
- Powerpoint Presentation: RNA-Sequencing, what does RNA-Seq data look like
- Pipeline for RNA-Seq alignment

Optional Exercise:

https://github.com/sabrsyed/InformaticsTools_2020/blob/master/01_Unix_QuickReview_ProblemSet.md

Why Cluster?

Massive data coming from Deep Sequencing needs to be

- stored
- (parallel) processed

It is not feasible to process this kind of data even using a high-end computer.

University of Massachusetts Green High Performance Computing Cluster

$\text{HPCC} \equiv \text{GHPCC} \equiv \text{MGHPCC} \equiv \text{the Cluster}$

HPC : High performance computing

Cluster : a number of similar things that occur together

Computer Cluster : A set of computers connected together
that work as a single unit

MGHPCC has over 10K+ cores available and 400+ TB of high performance storage. It is located in Holyoke MA and provides computing services to the five campuses of UMass.

cluster:126 [DokuWiki]


cluster:73 [DokuWiki]

cluster:136 [DokuWiki]

← → ↺

dokuwiki.wesleyan.edu/doku.php?id=cluster:126#description

☆ ⓘ 📄

DokuWiki

Log In

Search

Recent ChangesMedia ManagerSitemap

Trace: • [start](#) • **126**

Back

Brief Guide to HPCC

This page will be maintained and provide information to get users started using the compute cluster. It is a merger of the old “brief description” page and the “queue description” page.

Description

The High Performance Compute Cluster (HPCC) is comprised of several login nodes (all are on our domain *wesleyan.edu* behind VPN for off campus access)

- primary login node `cottontail` (Supermicro 4U), primary scheduler and snapshot engine for /home
- secondary login node `cottontail2` (HP Proliant G380 2U), backup scheduler
- secondary login node `swallowtail` (Dell PowerEdge 2950 2U), backup scheduler, databases
- sandbox `petaltail` (Dell PowerEdge 2950 2U), test box, Warewulf provisioning CentOS6
- sandbox `whitetail` (HP Proliant G380 2U), Warewulf OpenHPC provisioning CentOS7
- zenoss monitoring and alerting server `hpcmon` (supermicro 1U, centos6)
- NFS server `greentail52` (SuperMicro 36+2, 2U), /sanscratch
- (only log in when moving conternt) file server node `sharptail` (Supermicro 4U), /home NFS server
- DR node `sharptail2` (Supermicro 2U), disaster recovery for /home, off site (active users only)
- storage servers `rstore0` and `rstore2` (Supermicro 4U), NFS mounts and Samba shares (2x 120T)
- storage servers `rstore4` and `rstore6` (Supermicro 4U), NFS mounts and Samba shares (2x 220T)
- mindstore storage servers `mstore0/mstore1` (Supermicro 4U), available on HPC (2x 110T)

cluster:126

Table of Contents

- [Brief Guide to HPCC](#)
- [Description](#)
- [Our Queues](#)
- [Other Stuff](#)

All queues are available for job submissions via all login nodes. Some nodes on Infiniband switches for parallel computational jobs (queues: me256fd, hp12). Our total job slot count is roughly 2,144 with our physical core count 1,480. Our total teraflops compute capacity is about 58 cpu side, 25 gpu side (double precision floating point) and 702 gpu side (mixed mode). Our total memory footprint is about 528 GB gpu side, 8,532 GB cpu side.

Home directory file system are provided (via NFS or IPoIB) by the node `sharptail` (our file server) from a direct attached disk array. In total, 10 TB of /home disk space is accessible to the users. Node `greentail52` makes available 55 TB of scratch space at /sanscratch via NFS. In addition all nodes provide local scratch space at /localscratch (excludes queue tinymem). The scheduler automatically makes directories in both these scratch areas for each job (named after JOBPID). Backup services for /home are provided via disk-to-disk point-in-time snapshots from node `sharptail` to node `cottontail` disk arrays. (daily, weekly, monthly snapshots are mounted read only on `cottontail` for self-serve content retrievals). Some faculty have their home directories on node `ringtail` which provides 33 TB via /home33. Some faculty also have their own storage (2x 110 TB via /mindstore). In addition no-quota, no-backup user directories can be requested in /homeextra1 (7 T) or /homeextra2 (5 T). All home directories will migrate to a FreeNAS/ZFS appliance named `hpcstore` in 2020 (190T usable, scalable to 1.2P).

Two (old) Rstore storage servers each provide about 104 TB of usable backup space which is not mounted on the compute nodes. Each Rstore server's content is replicated to a dedicated passive standby server of same size, located in same data center but in different racks. As of Spring 2019 we have added two new Rstore servers of 220 T each, fully backed up with replication.

Our Queues

Commercial software has their own queue limited by available licenses. There are no scheduler license resources, just queue jobs up in appropriate queue. Commercial software jobs are processed on the nodes of mw256fd and mw128.

Queue	Nr Of Nodes	Total GB Mem Per Node	Total Cores In Queue	Switch	Hosts	Notes
stata	<i>na</i>	<i>na</i>	<i>na</i>	QDR Infiniband	<i>any host</i>	6 licenses

Note: Matlab and Mathematica now have "unlimited licenses".

Queue	Nr Of Nodes	Total GB Mem Per Node	Job Slots In Queue	Switch	Hosts	Notes
hp12	32	12	256	QDR infiniband	n1-n32	CPU
mwgpu	5	256	120	QDR infiniband	n33-n37	GPU & CPU
mw256fd	8	256	192	QDR infiniband	n38-n45	CPU
tinymem	14	32	448	gigabit ethernet	n39-n59	CPU
mw128	18	128	648	gigabit ethernet	n60-n77	CPU
amber128	1	128	24	gigabit ethernet	n78	GPU & CPU
exx96	12	96	432	gigabit ethernet	n79-n90	GPU & CPU

Reaching the Nodes

We do **NOT** use the head node (ghpcc06) to process big data.
We use the cluster nodes to process it.

How do we reach the nodes?

We submit our commands as jobs to a *job scheduler* and the job scheduler finds an available node for us having the sufficient resources (cores & memory.)

Job Scheduler

Job Scheduler is a software that manages the resources of a cluster system. It manages the program execution in the nodes. It puts the *jobs* in a (priority) queue and executes them on a node when the requested resources become available.

Let's submit another job and specify the resources this time.
To set


- 1 We explicitly state that we request a single core, `-n 1`
- 2 The memory limit to 1024 MB, we add `-R rusage[mem=1024]`
- 3 Time limit to 20 minutes, we add `-W 20`
- 4 Queue to short, we add `-q short`

```
$ bsub -n 1 -R rusage[mem=1024] -W 20 -q short "sleep 300"
```

We need 4 cores as we'll run our process in 4 threads, so we need `-n 4`.
2 GB = 2048 MB, so we need the parameter `-R rusage[mem=2048]`.
We can **estimate** the running time to be $20 / 4 = 5$ hours = 300 mins. So, let's ask for 330 mins to be on the safer side.

```
$ bsub -R span[hosts=1] -n 4 -R rusage[mem=2048] -W 330 -q long "~/bin/myscript.pl -p 4"
```

<https://dokuwiki.wesleyan.edu/doku.php?id=cluster:59>



DokuWiki

Log In

Search

[Recent Changes](#) [Media Manager](#) [Sitemap](#)

Trace: • [59](#)

[Back](#)

Complete Documentation

It's all at this link [COMPLETE DOCUMENTATION FOR LSF/HPC 6.2](#) and very good.

New Features in LSF 6.2

This page will be expanded to show examples of LSF/HPC advanced features.

The more information you can provide to the scheduler regarding run times, resources needed and when, the more efficient the scheduling will be. The examples below are just made up scenarios. Try to get familiar with them or ask for hands-on working sessions.

⇒ Also read up on the new queue configurations: [Link](#)

As part of the upgrade:

- Jobs were terminated ... for a list of which ones view [External Link](#)
- The working directories of those terminated jobs were saved in `/sanscratch/OLDJOBS`, help your self ...
- When the new scheduler came online it started with JOBPID 101 ... that may clobber some of your old output files so i've spooled the JOBPIDs forward to 30,000.

Table of Contents

- [Complete Documentation](#)
- [New Features in LSF 6.2](#)
 - [Exclusive](#)
 - [Resource Reservation](#)
 - [Wall Clock Time](#)
 - [Parallel Jobs](#)
 - [Old Way](#)
 - [Spanning](#)
 - [New Way](#)

cluster:59

Running Jobs on Linux/Cluster

http://barc.wi.mit.edu/education/hot_topics/lsf/Running_jobs_on_Linux_Cluster.pdf

<code>bjobs</code>	checking submitted jobs
<code>bjobs -a</code>	checking recently ended jobs
<code>bjobs -l JOBID</code>	see details of a particular job using the job id #
<code>bpeek JOBID</code>	peek at the stdout and stderr output of unfinished job
<code>bkill JOBID</code>	kills jobs

Quick Review

command	description
<code>ls</code>	list directory contents
<code>cd</code>	change directory
<code>mkdir</code>	make a directory
<code>rm</code>	remove, or delete files and directories. Use caution, it is easy to delete more that you want.
<code>head</code>	prints the top few lines to the terminal window
<code>tail</code>	prints the last few lines to the terminal window
<code>sort</code>	sorts the lines
<code>uniq</code>	prints the unique lines
<code>grep</code>	finds the lines that contain a pattern
<code>wc</code>	counts the number of lines, characters and words
<code>mv</code>	move files
<code>cp</code>	copy files
<code>date</code>	returns the current date and time
<code>pwd</code>	return working directory name
<code>ssh</code>	remote login
<code>scp</code>	remote secure copy
<code>~</code>	shortcut for your home directory
<code>man <command></code>	manual page for the command e.g. <code>man ls</code> to get the man page for <code>ls</code>
less or zless	read text files/read .gz compressed text files

Summer 2020

Unix/Linux for Informatic Analysis

https://github.com/sabrsyed/InformaticsTools_2020

Week 1

- Unix Primer for Biologists: Chapters U1 – U16
 - Learn how to use UNIX/Linux
- Logging in to the Cluster
 - Learn to navigate the Cluster
- Powerpoint Presentation: learn the technology behind genome sequencing, what does ChIP-Seq data look like
- Pipeline for ChIP alignment

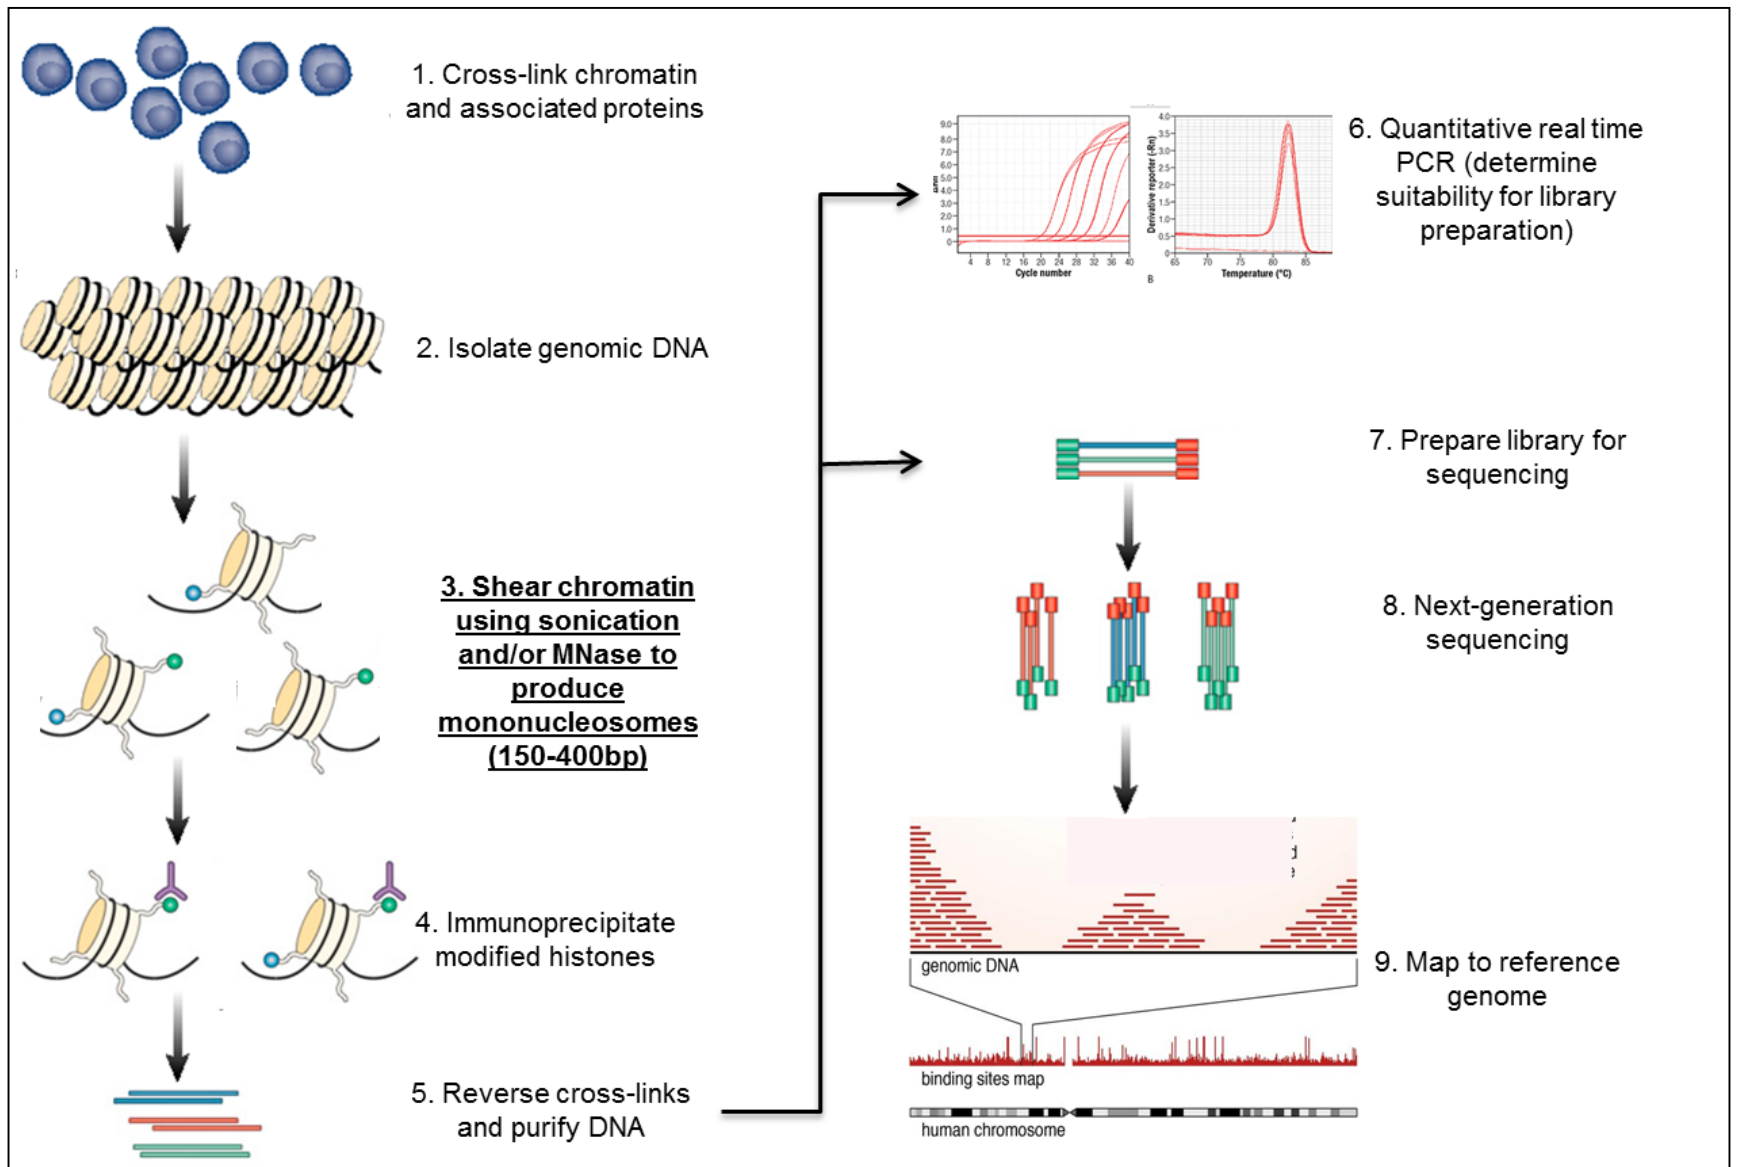
Week 2

- Unix Primer for Biologists: Chapters U17 – U34
 - Learn how to use UNIX/Linux
- Filezilla
 - Uploading files to the Cluster
- Powerpoint Presentation: RNA-Sequencing, what does RNA-Seq data look like
- Pipeline for RNA-Seq alignment


Optional Exercise:

https://github.com/sabrsyed/InformaticsTools_2020/blob/master/01_Unix_QuickReview_ProblemSet.md

ChIP-Seq (Chromatin Immunoprecipitation)



Control of gene expression by histone modifications



- The nucleosome is made up of dimers of core histones H2A, H2B, H3, H4 with 147 base pairs of double stranded DNA wrapped around the nucleosome



Activating histone modifications

H3K4me3
H3K27ac
H4ac
H3K4me1

Repressive histone modifications

H3K27me3
H3K9me3
H2AK119ubi1

ChIP-qPCR



**Activating
histone
modifications**

H3K4me3

H3K27ac

H4ac

**Repressive
histone
modifications**

H3K27me3

H3K9me3


H2AK119ubi1

H3K4me1

HiSeq 2000




FIGURE 1: ILLUMINA GENOME ANALYZER FLOW CELL




Up to eight samples can be loaded onto the flow cell for simultaneous analysis on the Illumina Genome Analyzer.

Illumina 2000 HiSeq


Bridge Amplification



Bridge Amplification



7. DETERMINE FIRST BASE





First chemistry cycle: to initiate the first sequencing cycle, add all four labeled reversible terminators, primers and DNA polymerase enzyme to the flow cell.




Raw Data is Images

- 8 channels per flow cell
- 300 tiles per channel
- 20,000 clusters/reads per tile

- 2400 images per reaction cycle
- 86,400 images for a 36-bp read length
- 700 GB of image data, 400 GB of files





Images




Image Analysis

Lane	File	X	Y	Spot1-FFFF	Spot2-FFFF
0	12	924	1000	955.9 388.9 369.7 228.4	955.9 922.9 388.9 367.9
0	12	973	980	955.9 915.9 2207.6 9199.9	955.9 999.9 915.9 999.9
0	12	985	790	955.9 790.9 45.9 47.4	955.9 999.9 45.9 47.7
0	12	990	990	955.9 990.9 95.9 795.9	955.9 995.7 995.9 999.9
0	12	1007	1007	955.9 915.9 957.9 919.9	955.9 995.9 915.9 919.9
0	12	1004	990	2207.6 990.9 47.2 45.1	955.9 915.9 957.9 919.9
0	12	987	990	955.9 990.9 42.2 385.9	955.9 999.9 42.2 385.9
0	12	997	1000	955.9 957.9 995.7 995.7	955.9 995.9 957.9 995.9
0	12	999	999	955.9 915.9 45.2 385.9	955.9 999.9 45.2 385.9
0	12	999	999	955.9 915.9 95.7 795.9	955.9 999.9 95.7 795.9
0	12	997	1000	955.9 990.9 955.9 955.9	955.9 995.9 955.9 955.9
0	12	997	1000	955.9 915.9 957.9 919.9	955.9 995.9 915.9 919.9

Base Calling

```

ATGGCCTGGGCTAGTTTCGATTACGA
CCTGGGCTAGTTTCGATTACGATCGA
GCTAGTTTCGATTACGATCGATCGTT
ATCGATCGTTGCATGCTGGGGTAGTG
TTCGATTACGATCGATCGTTGCATGCT
TCGATTACGATCGATCGTTGCATGCTG
CTAGTTTCGATTACGATCGATCGTTG
TCGATTACGATCGATCGTTGCATGCTG
TACGATCGATCGTTGCATGCTGGGGTA
TCGATCGTTGCATGCTGGGGTAGTGC
TCGATTACGATCGATCGTTGCATGCTG
CGATTACGATCGATCGTTGCATGCTG
TAGTTTCGATTACGATCGATCGTTGC
CGATTACGATCGATCGTTGCATGCTGG
ACGATCGATCGTTGCATGCTGGGGTAG
    
```

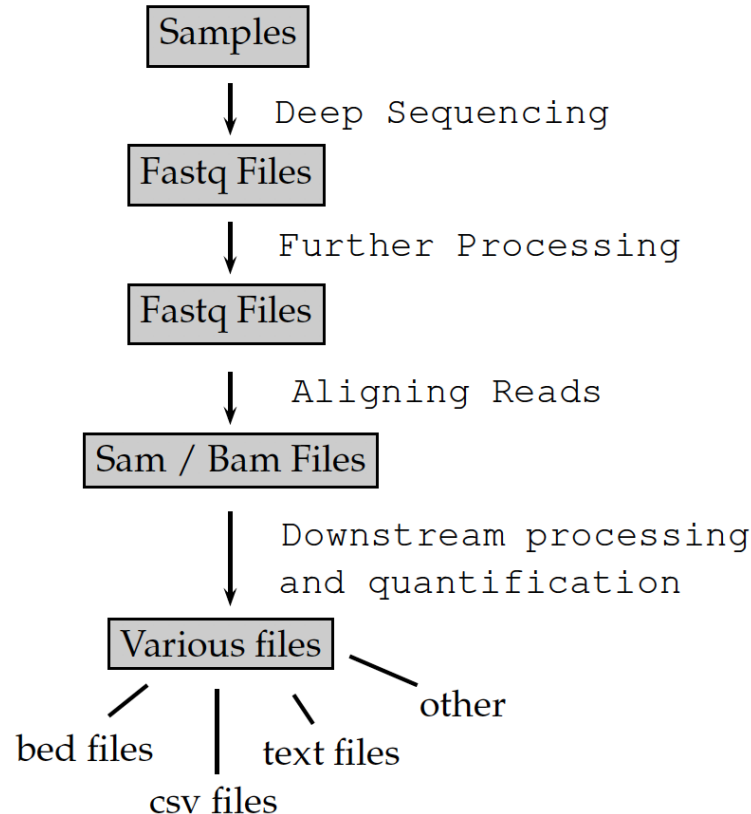
Aligned Reads

```

TGCCTAAGGCTAAGTTCATGCTAAGGTTTCGAA
A GCGTAAGGCTAAGTTCATGCTAAGGTTTCGAA
AT CGTAAGGCTAAGTTCATGCTAAGGTTTCGAA
ATG GTAAGGCTAAGTTCATGCTAAGGTTTCGAA
ATGC TAAGGCTAAGTTCATGCTAAGGTTTCGAA
ATCG AAGGCTAAGTTCATGCTAAGGTTTCGAA
ATGCGT AAGGCTAAGTTCATGCTAAGGTTTCGAA
ATGCGTA GCTAAGTTCATGCTAAGGTTTCGAA
ATGCGTAA CTAGTTCATGCTAAGGTTTCGAA
    
```

- Interpret images as intensities
- Convert intensities to base calls
- Assemble reads into complete sequence
 - ⇒ 36-bp read length
 - ⇒ 4 to 6 million reads per flow cell lane

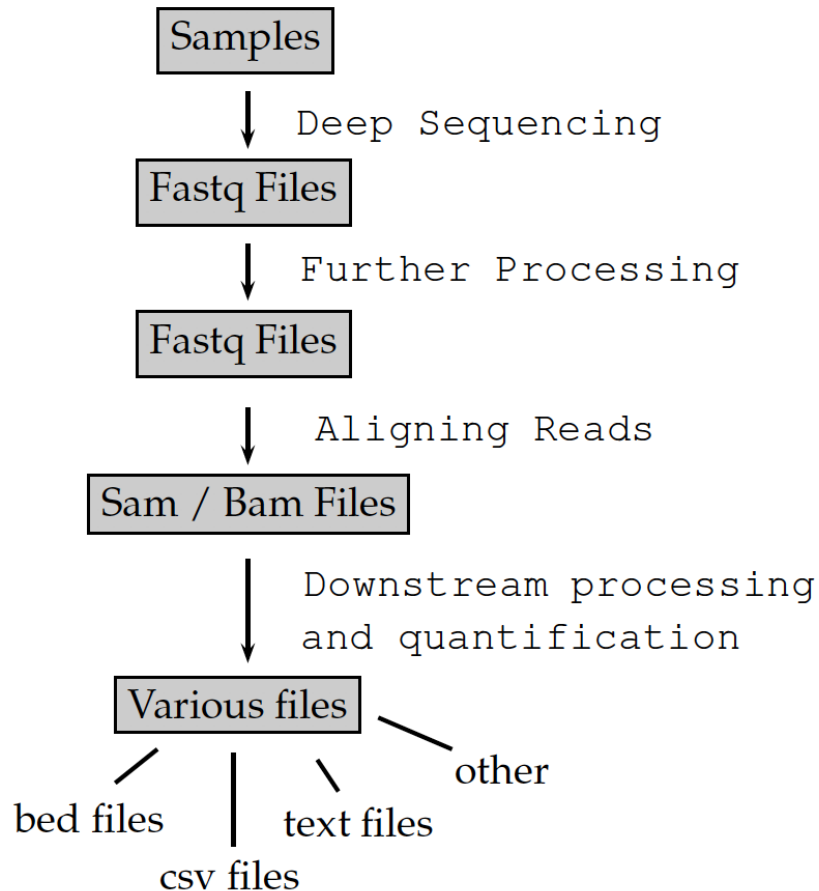
A Typical Deep-Sequencing Workflow



Deep Sequencing Data pipelines involve a lot of text processing.

This is an oversimplified model and your workflow can look different from this!

SAM / BAM Files



When a fastq file is aligned against a reference genome, a sam or a bam file is created as the ultimate output of the alignment. These files tell us where and how reads in the fastq file are mapped.

Sequencing Workflow

- **FASTQ**: a text-based format for storing nucleotide sequences (reads) and their quality scores.

[1]

```

CCCFFFFFFFFHHHJJJJJJJJIIJJJJJJJDHIJJDHIIJJJJJJJIIJJJJJG
@R0212989:323:C3P6FACXX:1:1211:5383:20897/1
CCACAGTGTACTTTATTTAATGATTTTTGTACTTTGTGTTGCAATAAAATA
+
CCCCFFFDHFFFFFFHJJJJJJJJJJIIJJJJGGHIGIJJIDFHIJJJIHJJJJJJJ
@R0212989:323:C3P6FACXX:1:1306:4306:19653/1
CAACTTGTAAGTGTGTCTTTCTTGTTGGAGGCTGCTGCCCTGGGCAGTGA
+
CCCCFFFFDHHGGFHHEHHIIJJIIIIJIHJJJGIJGGEH@GHGGGIJJGCGHIG
@R0212989:323:C3P6FACXX:1:1305:10910:89723/1
GCAAATACTCCACACACTGTGCTTTGAGCTAGAGCACTTGGAGTCACTGCC
+
CCCCFFFFFHHHHHJJJIIJJJJJJJJJJIIJJJJGJGGHIIIJJIEHHHJJJIH
@R0212989:323:C3P6FACXX:1:2316:1690:48422/1
CAGACCTTCCTTTAGAATTCAACTTGTAAGTGTGTCTTTCTTGTTGGAGG
+
@C@FDDFFFHDFFBFEGEGBFHIIJJGGJICHCFGGIIJJIGIG9GEHBGHG
@R0212989:323:C3P6FACXX:1:2108:10999:30758/1
CATAACCAGACCTTCCTTTAGAATTCAACTTGTAAGTGTGTCTTTCTTGTT
+
CCCCFFFFFHHHHHJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJ
@R0212989:323:C3P6FACXX:1:2211:12636:46495/1
TGGTCTGGTTATGTGGGGTTGGAATATGTATATCTATATATCTCTATATAT
+
@CCDDFFFFFHHHHHJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJ
@R0212989:323:C3P6FACXX:1:2106:20600:69959/1
TTTACATCAAAGAATTAACTCATTTAAATTTTTGTTTCAAATAAAAT
+
=@@D?DDDFFHHDBBFGEBBHFIIBHEHHG4<CFFHGHEHCFFGFIIIEG
@R0212989:323:C3P6FACXX:1:1212:8997:69630/1
CTGATGCATAGTCGGTGACATTCTTGAGTTTCTCTCTCCATTTCAGAAATA
+
CCCCFFFFFHHHHHJJJHIIJJJJJJJJJJJJHIIJJJIHIIIIJJJIIJGJJJJ
@R0212989:323:C3P6FACXX:1:1106:5824:44089/1
TATATCTTCACGTTGCCTGCACACACCTTATTTCTGAAATGGAGAGAGAAAA

```

Sequencing Workflow

- BAM: The Sequence Alignment/Mapping (SAM) format is a text-based format for storing read alignments against reference sequences and it is interconvertible with the binary BAM format.

Each alignment line has 11 mandatory fields for essential alignment information such as mapping position, and variable number of optional fields for flexible or aligner specific information.

```
@HD VN:1.5 SO:coordinate
@SQ SN:ref LN:45
r001 99 ref 7 30 8M2I4M1D3M = 37 39 TTAGATAAAGGATACTG *
r002 0 ref 9 30 3S6M1P1I4M * 0 0 AAAAGATAAGGATA *
r003 0 ref 9 30 5S6M * 0 0 GCCTAAGCTAA * SA:Z:ref,29,-,6H5M,17,0;
```


Sequencing Workflow

•BED file


chr1	3001975	3002012	-	SL-XAR_1_FC13498AAXX_6_127_905_305	0
chr1	3004386	3004423	-	SL-XAR_1_FC13498AAXX_6_219_329_203	0
chr1	3030430	3030467	+	SL-XAR_1_FC13498AAXX_6_132_674_273	0
chr1	3031032	3031069	-	SL-XAR_1_FC13498AAXX_6_127_680_131	0
chr1	3033263	3033300	+	SL-XAR_1_FC13498AAXX_6_187_430_40	1
chr1	3035898	3035935	+	SL-XAR_1_FC13498AAXX_6_137_684_268	0
chr1	3036679	3036716	-	SL-XAR_1_FC13498AAXX_6_155_848_458	1
chr1	3043625	3043662	-	SL-XAR_1_FC13498AAXX_6_180_603_902	0
chr1	3044153	3044190	-	SL-XAR_1_FC13498AAXX_6_197_866_908	0
chr1	3044528	3044565	-	SL-XAR_1_FC13498AAXX_6_202_521_367	1
chr1	3045627	3045664	-	SL-XAR_1_FC13498AAXX_6_153_959_874	0
chr1	3053181	3053218	+	SL-XAR_1_FC13498AAXX_6_183_138_309	0
chr1	3062755	3062792	+	SL-XAR_1_FC13498AAXX_6_178_383_87	1
chr1	3065421	3065458	-	SL-XAR_1_FC13498AAXX_6_205_876_214	0
chr1	3066969	3067006	+	SL-XAR_1_FC13498AAXX_6_171_205_595	1
chr1	3067298	3067335	+	SL-XAR_1_FC13498AAXX_6_213_767_278	0
chr1	3067600	3067637	-	SL-XAR_1_FC13498AAXX_6_202_441_205	1
chr1	3067721	3067758	+	SL-XAR_1_FC13498AAXX_6_144_842_179	0
chr1	3073695	3073732	-	SL-XAR_1_FC13498AAXX_6_128_477_62	6
chr1	3080674	3080711	+	SL-XAR_1_FC13498AAXX_6_134_539_5	0
chr1	3082545	3082582	-	SL-XAR_1_FC13498AAXX_6_191_700_362	0
chr1	3082596	3082633	-	SL-XAR_1_FC13498AAXX_6_171_157_717	4
chr1	3090549	3090586	+	SL-XAR_1_FC13498AAXX_6_144_874_773	0
chr1	3094861	3094898	+	SL-XAR_1_FC13498AAXX_6_189_889_433	1
chr1	3097811	3097848	-	SL-XAR_1_FC13498AAXX_6_193_600_983	0
chr1	3098704	3098741	+	SL-XAR_1_FC13498AAXX_6_133_38_81	0

Binding profile

1 kb



50 kb



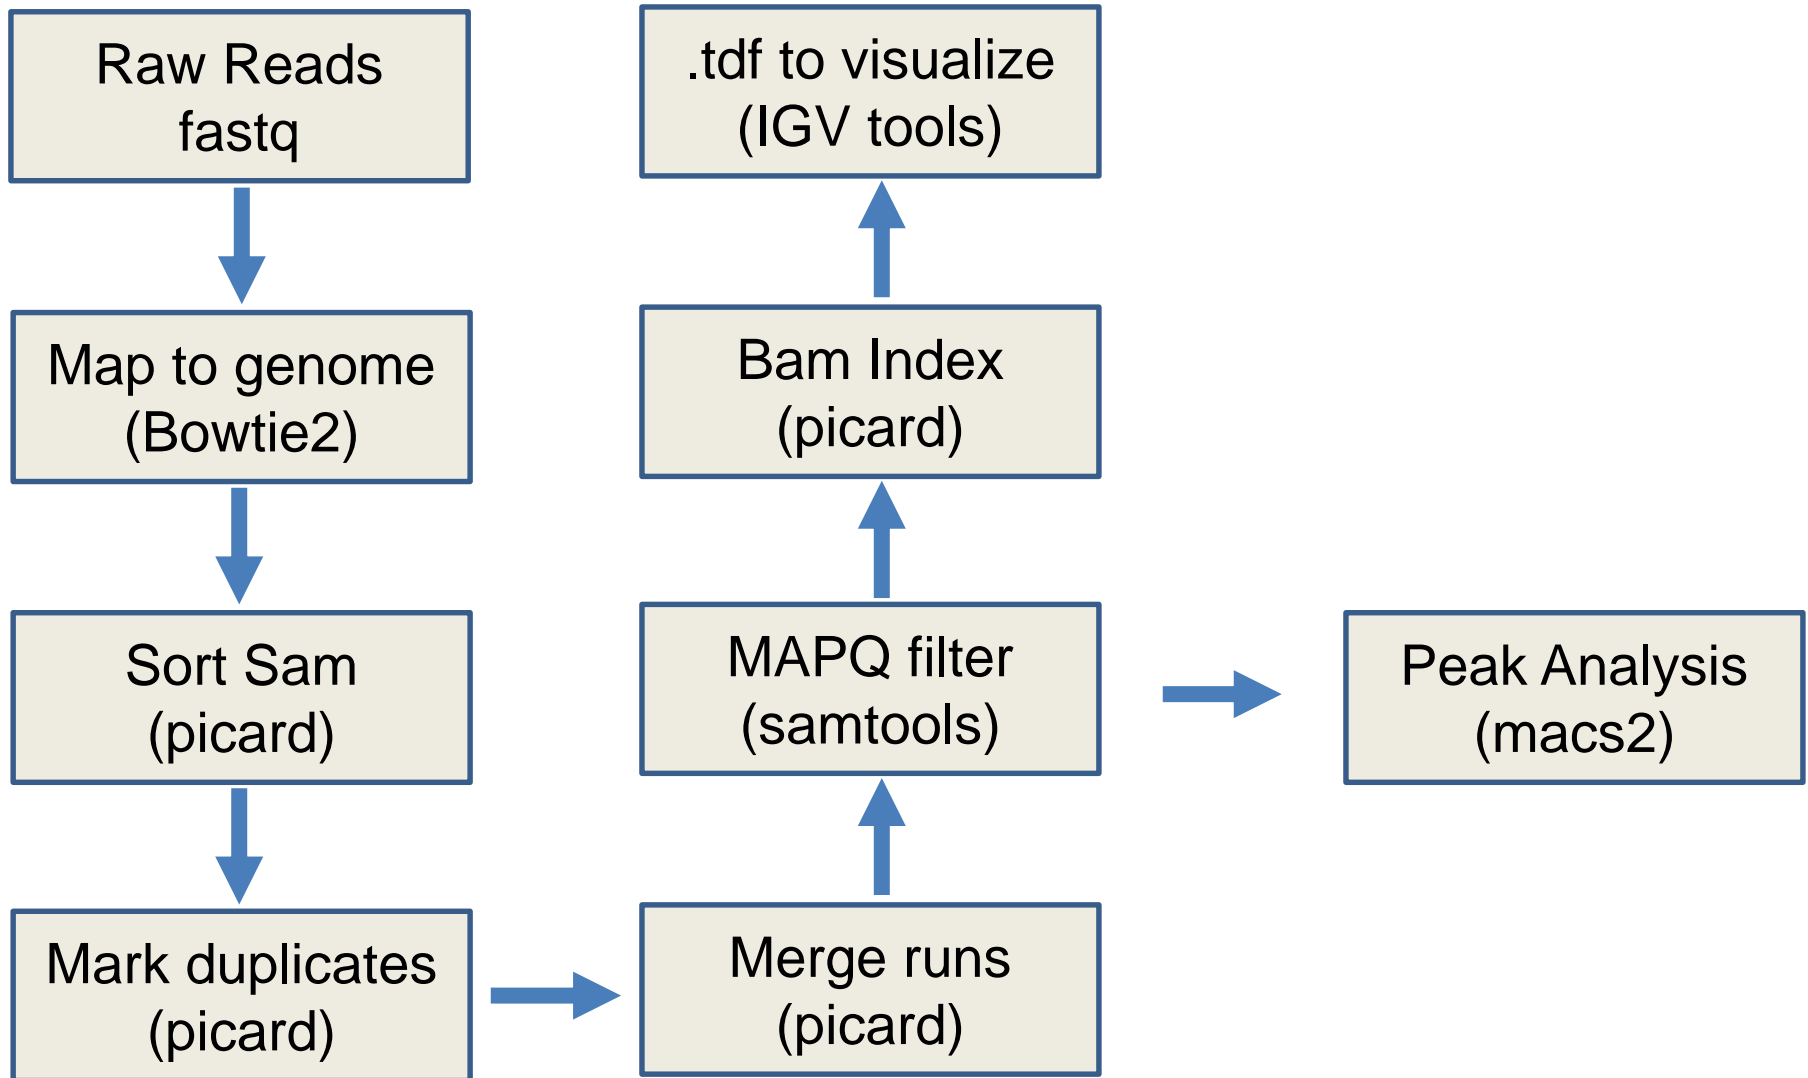
1 kb

RNAP II
ZFP36
(human Chr 19)



Kagey et al. (2010). Nature 467: 430
Mikkelsen et al. (2007). Nature 448: 553
Pepke et al. (2009). Nat. Methods 6: S22

My ChIP-Sequencing Workflow



Summer 2020

Unix/Linux for Informatic Analysis

https://github.com/sabrsyed/InformaticsTools_2020

Week 1

- Unix Primer for Biologists: Chapters U1 – U16
 - Learn how to use UNIX/Linux
- Logging in to the Cluster
 - Learn to navigate the Cluster
- Powerpoint Presentation: learn the technology behind genome sequencing, what does ChIP-Seq data look like
- Pipeline for ChIP alignment

Week 2

- Unix Primer for Biologists: Chapters U17 – U34
 - Learn how to use UNIX/Linux
- Filezilla
 - Uploading files to the Cluster
- Powerpoint Presentation: RNA-Sequencing, what does RNA-Seq data look like
- Pipeline for RNA-Seq alignment

Optional Exercise:

https://github.com/sabrsyed/InformaticsTools_2020/blob/master/01_Unix_QuickReview_ProblemSet.md